

Building a Linux Container using Namespace

Sumit Yadav

Overview

Ever wondered how Linux Containers work?

1. Docker is one of the most popular container implementations.
2. Containers share the same OS kernel but isolate application processes.
3. Key technology: Namespaces.

Namespaces in Linux Containers

Namespaces abstract global system resources to create isolated instances for processes.

IPC Namespace

- Isolates interprocess communication resources: Message Queues, Semaphores, and Shared Memory.

Network Namespace

- Isolates networking resources: Network devices, IP addresses, IP routing tables, /proc/net directory, and port numbers.

Mount Namespace

- Isolates filesystem mount points, allowing different views of the

PID Namespace

- Isolates the process ID number space, enabling the same PID in different namespaces.

User Namespace

- Isolates user and group ID number spaces, allowing different user/group IDs inside and outside the namespace.

UTS Namespace

- Isolates system identifiers: hostname and NIS domain name, allowing each container to have its own.

Installation and Setup

Every namespace is implemented using the `unshare` Linux command.
We will implement, build, and execute a simple container using `golang`.

Creating a Simple Container

Using `unshare`

```
sudo unshare -u /bin/sh
```

Options:

<code>-m, --mount[=<file>]</code>	unshare mounts namespace
<code>-u, --uts[=<file>]</code>	unshare UTS namespace (hostname etc)
<code>-i, --ipc[=<file>]</code>	unshare System V IPC namespace
<code>-n, --net[=<file>]</code>	unshare network namespace
<code>-p, --pid[=<file>]</code>	unshare pid namespace
<code>-U, --user[=<file>]</code>	unshare user namespace
<code>-C, --cgroup[=<file>]</code>	unshare cgroup namespace
<code>-f, --fork</code>	fork before launching <program>
<code>--mount-proc[=<dir>]</code>	mount proc filesystem first (implies <code>--mount</code>)
<code>-r, --map-root-user</code>	map current user to root (implies <code>--user</code>)
<code>--propagation</code>	slave shared private unchanged

Using `golang`

```
func execContainerShell() {
    log.Printf("Ready to exec container shell ...\n")

    if err := syscall.Sethostname([]byte("leopard")); err != nil {
        panic(err)
    }

    const sh = "/bin/sh"

    env := os.Environ()
    env = append(env, "PS1=-> ")

    if err := syscall.Exec(sh, []string{"", env}); err != nil {
        panic(err)
    }
}
```


UTS Namespace

Listing Namespaces

Each process has associated namespaces listed in `/proc/[PID]/ns`.
Running `ls -l /proc/$$/ns` shows symbolic links representing the namespaces of the current process.

Isolating Hostname with UTS Namespace

```
sudo unshare -u /bin/sh  
hostname leopard
```

This isolates the `hostname` in a new UTS namespace. Changes to the hostname in this namespace do not affect the parent namespace.

Go Program for Namespace Isolation

```
cmd.SysProcAttr = &syscall.SysProcAttr{  
    Cloneflags: syscall.CLONE_NEWUTS,  
}
```

- **Cloneflags:** Specifies the namespaces to unshare (UTS).

User Namespace

Launching an Isolated Shell

```
sudo unshare -uU /bin/sh  
>> uid=65534(nobody) gid=65534(nogroup) groups=65534(nogroup)
```

Mapping the user/group ID to the parent user/group ID:

```
sudo unshare -uUr /bin/sh  
>> uid=65534(root) gid=65534(root) groups=65534(root)
```

Using `golang`

```
cmd.SysProcAttr = &syscall.SysProcAttr{
    Cloneflags: syscall.CLONE_NEWUTS | syscall.CLONE_NEWUSER,
    UidMappings: []syscall.SysProcIDMap{
        {ContainerID: 0, HostID: 0, Size: 1},
    },
    GidMappings: []syscall.SysProcIDMap{
        {ContainerID: 0, HostID: 0, Size: 1},
    },
}
```

- **Cloneflags:** Specifies the namespaces to unshare (UTS and user).
- **UidMappings** and **GidMappings:** Maps the root user (ID 0) in the container to the root user on the host.

PID Namespace

Launching an Isolated Shell

```
sudo unshare -uUrpf --mount-proc /bin/sh  
ps -fu
```

This shows the isolation between the new namespace and the parent namespace.

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	4628	880	pts/1	S	09:08	0:00	/bin/sh
root	6	0.0	0.0	37368	3340	pts/1	R+	09:12	0:00	ps -fu

Using `golang`

```
cmd.SysProcAttr = &syscall.SysProcAttr{
    Cloneflags: syscall.CLONE_NEWUTS | syscall.CLONE_NEWUSER | syscall.CLONE_NEWNS | syscall.CLONE_NEWPID,
    UidMappings: []syscall.SysProcIDMap{
        {ContainerID: 0, HostID: 0, Size: 1},
    },
    GidMappings: []syscall.SysProcIDMap{
        {ContainerID: 0, HostID: 0, Size: 1},
    },
}
```

- We specify the additional `syscall.CLONE_NEWNS` and `syscall.CLONE_NEWPID` OS attributes to indicate the command be run in a new PID namespace.

Mount Namespace

Steps to Set Up and Isolate a Minimal Ubuntu Base Image

1. Prepare the Base Image:

- Create necessary directories: `mkdir -p /tmp/rootfs/.old_root`
- Extract the Ubuntu base image:
`tar -xvf $HOME/Downloads/ubuntu-base-18.04.4-base-amd64.tar.gz --directory /tmp/rootfs`
- Navigate to `/tmp` directory.

2. Isolate the Environment with Unshare:

- Launch a container with isolated UTS, User, PID, and Mount namespaces: `sudo unshare -uUrpfm --mount-proc /bin/sh`

3. View and Compare Mount Points:

- In the parent namespace, list all mount points:

```
cat /proc/mounts | sort
```

- In the new namespace, list all mount points:

```
cat /proc/mounts | sort
```

- Compare the mount points to observe isolation.

4. Modify the New Namespace:

- Make the root filesystem private: `mount --make-rprivate /`
- Bind mount the root filesystem: `mount --rbind rootfs/ rootfs/`
- Mount the proc filesystem: `mount -t proc proc rootfs/proc`
- Change the root filesystem using pivot_root:
`pivot_root rootfs/ rootfs/.old_root`
- Change to the new root directory: `cd /`

5. Validate Changes and Clean Up:

- Verify the isolated environment by creating and checking files in `/tmp`.
- Mount `/tmp` as tmpfs: `mount -t tmpfs tmpfs /tmp`
- Create a file in the new namespace:
`echo 'leopard' > /tmp/leopard.txt`
- List the file in both namespaces to confirm isolation.
- Remove the parent root filesystem:
`mount --make-rprivate /.old_root` and `umount -l /.old_root`
- Exit the new namespace with `exit`.

Using `golang`

1. Initial Setup:

- The program starts by checking its arguments. If the argument `"CLONE"` is passed, it calls `execContainerShell`. Otherwise, it forks a new process with the necessary namespace flags.

2. Creating Namespaces:

- A new process is forked with `CLONE_NEWUTS`, `CLONE_NEWUSER`, `CLONE_NEWNS`, and `CLONE_NEWPID` flags, isolating UTS, User, Mount, and PID namespaces.
- User and group ID mappings are set to map the root user inside the container to the root user on the host.

- Changes the hostname to `"leopard"`.
- Changes the current directory to `/tmp`.
- Makes the root filesystem private and recursively bind mounts the `rootfs/` directory.
- Uses `pivot_root` to change the root filesystem to `rootfs/`, effectively isolating it from the parent filesystem.
- Mounts a new `tmpfs` at `/tmp` and the `proc` filesystem at `/proc`.
- Calls `createTxtFile` to create

a file `/tmp/leopard.txt` with the content "leopard".

- Makes `.old_root` private and unmounts it to finalize the filesystem isolation.
- Replaces the current process with a new shell (`/bin/sh`) with a

4. Main Function Workflow:

- Checks for the `"CLONE"` argument and decides whether to create namespaces or execute the container setup.
- If namespaces are to be created, it sets up command attributes and runs a new instance of the program in the isolated namespaces.

Network Namespace

Launching an Isolated Shell

```
sudo unshare -uUrpfmn --mount-proc /bin/sh
ip link
ip link set dev lo up
ping 127.0.0.1 -c3
```

Host Terminal Setup

```
sudo brctl addbr br0
sudo brctl show
sudo ip addr add 172.20.1.2/24 dev br0
sudo ip link set br0 up
ip link
```

Connecting veth Pairs

Host Terminal

```
sudo ip link add veth0 type veth peer name veth1
export UPID=$(pidof unshare)
sudo ip link set veth1 netns $UPID
ip link
```

Namespace Terminal

```
ip link
```

Configuring IPs

Host Terminal

```
sudo ip addr add 172.20.1.3/24 dev veth0
```

```
sudo ip link set veth0 up
```

```
ip addr add 172.20.1.4/24 dev veth1
```

Namespace Terminal

```
ip link set veth1 up
```

```
ping 172.20.1.4 -c3
```

```
ping 172.20.1.3 -c3
```

Using `golang`

The Go program creates a container-like environment with network namespace isolation using a bridge and virtual Ethernet (veth) pairs.

1. Setup Bridge and Veth Pairs:

- Creates a bridge `br0` and assigns it an IP address.
- Creates a veth pair (`veth0`, `veth1`), attaches `veth0` to `br0`, and configures their IPs.

2. Container Initialization (`execContainerShell`):

- Sets the hostname, changes the working directory, and mounts necessary filesystems.
- Configures network interfaces (`lo` and `veth1`) inside the container.
- Runs a shell (`/bin/sh`) in the isolated environment.

3. Main Function:

- If the `"CLONE"` argument is provided, it enters the container initialization.
- Otherwise, it sets up the bridge and veth pairs, starts a new process with namespace isolation, and moves `veth1` to the new process's network namespace.
- Cleans up by deleting the bridge after the container process exits.

This program effectively isolates processes with separate networking setups, mimicking basic container behavior.

Thanks <3